

1. Write a CUDA program for to multiply each element of vector by 5 using multiple blocks and multiple threads.

```
#include<iostream>

using namespace std;

__global__ void mul(int *a, int *b, int *c){
    int j = blockDim.x;
    // blockDim specifies no. of threads in each block
    int i = blockIdx.x*j + threadIdx.x;
    c[i] = b[i]*a[i];
    // c[i] = i;
}nv

int main(){
    int a[6],b[6],c[6];
    for(int i=0; i<6; i++){
        a[i] = 2*i+11;
        b[i] = 4*i+7;
    }
    int *da, *db, *dc;
    cudaMalloc(&da, 6*sizeof(int));
    cudaMalloc(&db, 6*sizeof(int));
    cudaMalloc(&dc, 6*sizeof(int));
    cudaMemcpy(da, &a, 6*sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(db, &b, 6*sizeof(int), cudaMemcpyHostToDevice);
    mul<<<2,3>>>(da,db,dc);
    cudaMemcpy(&c, dc, 6*sizeof(int), cudaMemcpyDeviceToHost);
    for (int j=0; j<6; j++){
        cout<<b[j]<<" * "<<a[j]<<" = "<<c[j]<<endl;
    }
    cudaFree(da);
    cudaFree(db);
    cudaFree(dc);
    return 0;
}
```

2. Write a CUDA program for pairwise sum of elements of vector to showcase concept of syncthreads.

```
#include<iostream>

using namespace std;

__global__ void fun(int *a, int *b){
    int t = threadIdx.x;
    int n = blockDim.x;
    while(n!=0){
        if (t<n){
```

```

        // eg. a[0] += a[0+n], similiary for other indices, this
        would resuse the array again and again and keep on adding values.
        a[t] += a[t+n];
    }
    __syncthreads();
    n = n/2;
}
*b = a[0];
}
int main(){
    int N = 8;
    int a[N], b;
    for(int i=0; i<N; i++){
        a[i] = 2*i+11;
    }
    int *da, *db;
    cudaMalloc(&da, N*sizeof(int));
    cudaMalloc(&db, sizeof(int));
    cudaMemcpy(da, &a, N*sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(db, &b, sizeof(int), cudaMemcpyHostToDevice);
    fun<<<1,N/2>>>>(da, db);
    cudaMemcpy(&b, db, sizeof(int), cudaMemcpyDeviceToHost);
    cout<<"Res: "<<b<<endl;
    cudaFree(da);
    cudaFree(db);
    return 0;
}

```

3. Write a CUDA program for dot product using one block to showcase concept of shared memory.

```

#include<iostream>

using namespace std;
__global__ void dot_product(int *a, int *b, int *c){
    int i = threadIdx.x;
    // this allows accessing shared memory of all the threads
    of a block
    __shared__ int temp[6];
    temp[i] = b[i] * a[i];
    // this will ensure completion of all threads
    __syncthreads();
    if (threadIdx.x == 0){
        int res = 0;
        for (int i=0; i<6; i++){
            res += temp[i];
        }
    }
}

```

```
    }
    *c = res;
}

}

int main(){
    int size = 6;
    int a[size],b[size],c;
    cout<<"Enter elements of a: ";
    for(int i=0; i<size; i++){
        cin>>a[i];
    }
    cout<<"Enter elements of b: ";
    for(int i=0; i<size; i++){
        cin>>b[i];
    }

    int *da, *db, *dc;
    cudaMalloc(&da, size*sizeof(int));
    cudaMalloc(&db, size*sizeof(int));
    cudaMalloc(&dc, sizeof(int)
        cudaMemcpy(da, &a, size*sizeof(int),
cudaMemcpyHostToDevice);
        cudaMemcpy(db, &b, size*sizeof(int),
cudaMemcpyHostToDevice);
        dot_product<<<1,6>>>(da,db,dc);
        cudaMemcpy(&c, dc, sizeof(int), cudaMemcpyDeviceToHost);
        cout<<c<<endl;
        cudaFree(da);
        cudaFree(db);
        cudaFree(dc);
        return 0;
}
```